

CMPT 423/820 Question 6

- Your Name
- Your student number
- Your NSID

In this question, we'll get acquainted with Pandas DataFrames, methods, indexing, and elementary data visualization.

```
In [1]: # this is the conventional import
import pandas as pd
import matplotlib.pyplot as plt
```

Task 1

There's a datafile named `iris.csv` included with the assignment material. It's a classical dataset, very well-known (you can Google more information about it if you like -- we're using it to practice with the Pandas module). The file has 5 columns:

- Columns 1-4 are numerical measurements
- Column 5 is label, in string form. There are exactly 3 labels: `setosa`, `versicolor`, `virginica`. This label would be the output for a classifier built from this data.

Use Pandas function `read_csv()` to load the dataset.

```
In [2]: dataframe = pd.read_csv('iris.csv',
                                header=None,
                                names=[ 'SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm', 'Species' ],
                                index_col=False)
```

Task 2

The `DataFrame` class has a vast number of methods and functions, which could be of use in one application or another. The easiest ones summarize or visualize the data contained in the dataset.

In the cell below, use the `describe()` method to display some summary statistics.

```
In [3]: dataframe.describe()
```

Out[3]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

Task 3

In the cell below, use the `cov()` method to display some more summary statistics. What do you think this method does?

```
In [4]: dataframe.cov()
```

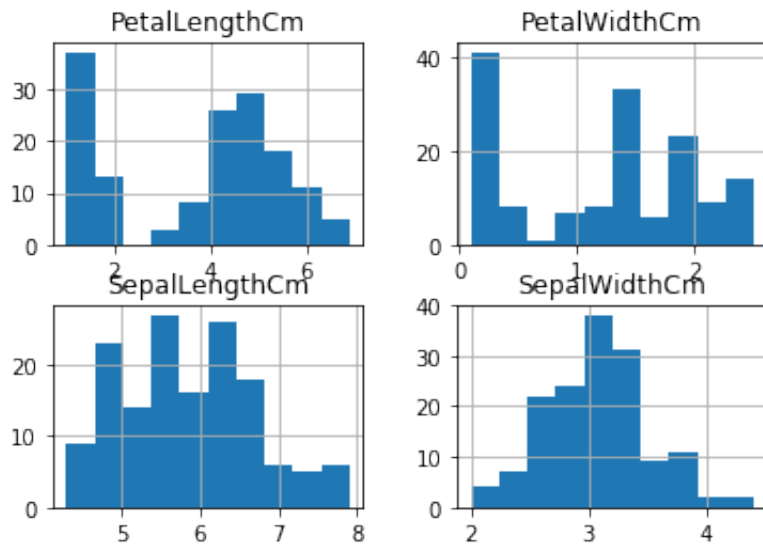
Out[4]:

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
SepalLengthCm	0.685694	-0.039268	1.273682	0.516904
SepalWidthCm	-0.039268	0.188004	-0.321713	-0.117981
PetalLengthCm	1.273682	-0.321713	3.113179	1.296387
PetalWidthCm	0.516904	-0.117981	1.296387	0.582414

Task 4

Summary statistics are informative, but a picture would be nice. Pandas has another simple method to help visualize the data, called `hist()`. Use this method to display histograms for the data.

```
In [5]: dataframe.hist()
plt.show()
```

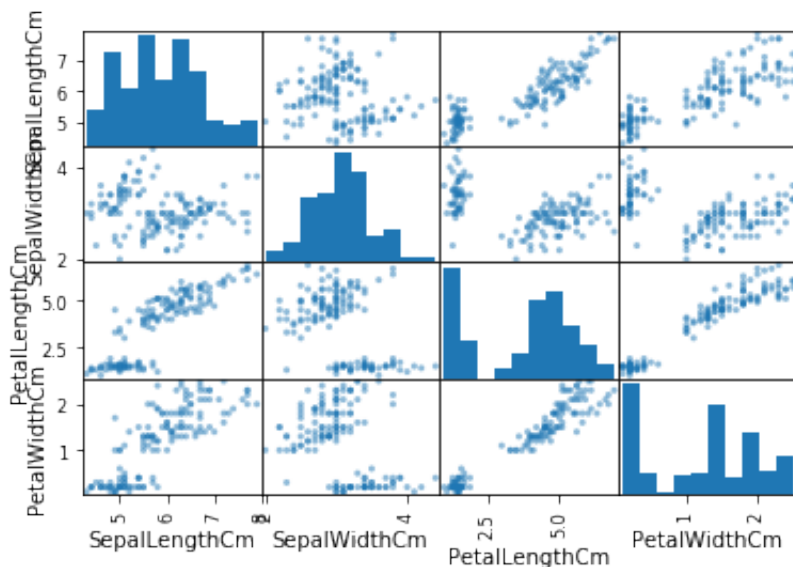


Task 5

In Task 3, we asked Pandas to display covariance information. Sometimes it's valuable to view data showing pair-wise scatter plots. Use Pandas' function `scatter_matrix()` to display the dataframe.

Hint: Notice the histograms along the diagonal!

```
In [6]: pd.plotting.scatter_matrix(dataframe)
plt.show()
```



What to hand in

Your version of this notebook named A1Q6.pdf, containing completed work above, and your name and student number at the top.

Evaluation:

- 1 mark. For Task 1, you used `read_csv()` to load a datafile into the notebook.
- 1 mark. For Task 2, you used `describe()` to display some information about the DataFrames.
- 1 mark. For Task 3, you used `cov()` to display some covariance information about the dataframe.
- 1 mark. For Task 4, you used `hist()` to display histograms for the dataframe.
- 1 mark. For Task 5, you used `scatter_matrix()` to display visualization of the original dataframe.

Grading:

- 1 mark. For Task 1, you used `read_csv()` to load a datafile into the notebook.
- 1 mark. For Task 2, you used `describe()` to display some information about the DataFrames.
- 1 mark. For Task 3, you used `cov()` to display some covariance information about the dataframe.
- 1 mark. For Task 4, you used `hist()` to display histograms for the dataframe.
- 1 mark. For Task 5, you used `scatter_matrix()` to display visualization of the original dataframe.
- **Deductions:**
 - If the required function was not used, dont give credit for that part.